全国大学生数学建模竞赛

厦门大学

学　　生 杨 成　王根盛　李嘉乐

指导教师 谭 忠

在2024年全国大学生数学建模竞赛中获得福建赛区本科组一等奖。

福建省大学生数学建模竞赛组委会

二〇二四年十二月

MIHBench: Benchmarking and Mitigating Multi-Image Hallucinations in Multimodal Large Language Models

ACM Multimedia 2025

Your Submissions | Author Tasks

#	Submission Summary	Official Review	Decision
1569	**MIHBench: Benchmarking and Mitigating Multi-Image Hallucinations in Multimodal Large Language Models** Download PDF Jiale Li, Mingrui Wu, Zixiang Jin, Hao Chen, Jiayi Ji, Xiaoshuai Sun, Liujuan Cao, Rongrong Ji ACMMM 2025 Show details	**4 Official Reviews Submitted** **Reviewe** ng: 7 / Confidence: 4 Read Official Review **Review** Rating: 4 / Confidence: 4 Read Official Review **Revi** Rating: 7 / Confidence: 4 Read Official Review **Revie** Rating: 5 / Confidence: 4 Read Official Review **Average Rating:** 5.75 (Min: 4, Max: 7) **Average Confidence:** 4 (Min: 4, Max: 4)	**ACMMM 2025 Recommendation:** **Accept** Read

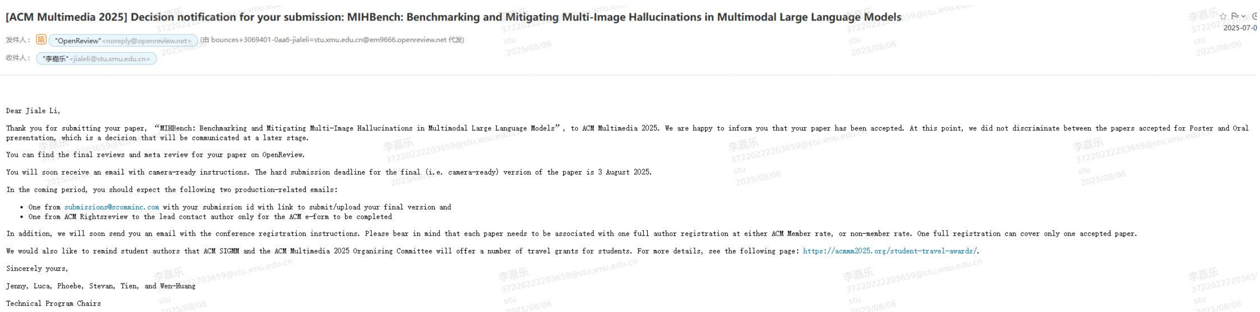

[ACM Multimedia 2025] Decision notification for your submission: MIHBench: Benchmarking and Mitigating Multi-Image Hallucinations in Multimodal Large Language Models

2025-07-0

发件人： "OpenReview" <noreply@openreview.net> (由 bounces+3069401-0aa6-jialeli=stu.xmu.edu.cn@em9666.openreview.net 代发)

收件人： "李嘉乐" <jialeli@stu.xmu.edu.cn>

Dear Jiale Li,

Thank you for submitting your paper, "MIHBench: Benchmarking and Mitigating Multi-Image Hallucinations in Multimodal Large Language Models", to ACM Multimedia 2025. We are happy to inform you that your paper has been accepted. At this point, we did not discriminate between the papers accepted for Poster and Oral presentation, which is a decision that will be communicated at a later stage.

You can find the final reviews and meta review for your paper on OpenReview.

You will soon receive an email with camera-ready instructions. The hard submission deadline for the final (i.e. camera-ready) version of the paper is 3 August 2025.

In the coming period, you should expect the following two production-related emails:

- One from submissions@scomminc.com with your submission id with link to submit/upload your final version and
- One from ACM Rightsreview to the lead contact author only for the ACM e-form to be completed

In addition, we will soon send you an email with the conference registration instructions. Please bear in mind that each paper needs to be associated with one full author registration at either ACM Member rate, or non-member rate. One full registration can cover only one accepted paper.

We would also like to remind student authors that ACM SIGMM and the ACM Multimedia 2025 Organising Committee will offer a number of travel grants for students. For more details, see the following page: https://acmmm2025.org/student-travel-awards/.

Sincerely yours,

Jenny, Luca, Phoebe, Stevan, Tien, and Wen-Huang

Technical Program Chairs

MIHBench: Benchmarking and Mitigating Multi-Image Hallucinations in Multimodal Large Language Models

Jiale Li*
jialeli@stu.xmu.edu.cn
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

Mingrui Wu*
mingrui0001@gmail.com
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China
Zhongguancun Academy
Beijing, China

Zixiang Jin
37220222203643@stu.xmu.edu.cn
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

Hao Chen
37120222203278@stu.xmu.edu.cn
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

Jiayi Ji†
jjyxmu@gmail.com
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

Xiaoshuai Sun
xssun@xmu.edu.cn
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

Liujuan Cao
caoliujuan@xmu.edu.cn
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

Rongrong Ji
rrji@xmu.edu.cn
Xiamen University
Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China
Xiamen, Fujian, China

ABSTRACT

Despite growing interest in hallucination in Multimodal Large Language Models (MLLMs), existing studies primarily focus on single-image settings, leaving hallucination in multi-image scenarios largely unexplored. To address this gap, we conduct the first systematic study of hallucinations in multi-image MLLMs and propose **MIHBench**, a benchmark specifically tailored for evaluating object-related hallucinations across multiple images. MIHBench comprises three core tasks—Multi-Image Object Existence Hallucination, Multi-Image Object Count Hallucination, and Object Identity Consistency Hallucination—targeting semantic understanding across object existence, quantity reasoning, and cross-view identity consistency. Through extensive evaluation, we identify key factors associated with the occurrence of multi-image hallucinations, including: (1) a progressive relationship between the number of image inputs and the likelihood of hallucination occurrences; (2) a strong correlation between single-image hallucination tendencies and those observed in multi-image contexts; and (3) the influence of same object image ratios and the positional placement of negative samples within image sequences on the occurrence of object identity consistency hallucination. To address these challenges, we propose a **Dynamic Attention Balancing (DAB)** mechanism that adjusts inter-image attention distributions while preserving the overall visual attention proportion. Experiments across multiple state-of-the-art MLLMs demonstrate that our method effectively reduces hallucination occurrences and enhances semantic integration and reasoning stability in multi-image scenarios. The project page is available at https://github.com/pgtrece/DAB/.

*Equal contribution.
†Corresponding author.

MM '25, October 27–31, 2025, Dublin, Ireland

ACM ISBN 979-8-4007-2035-2/2025/10
https://doi.org/10.1145/3746027.3754993

ControlMLLM: Training-Free Visual Prompt Learning for Multimodal Large Language Models

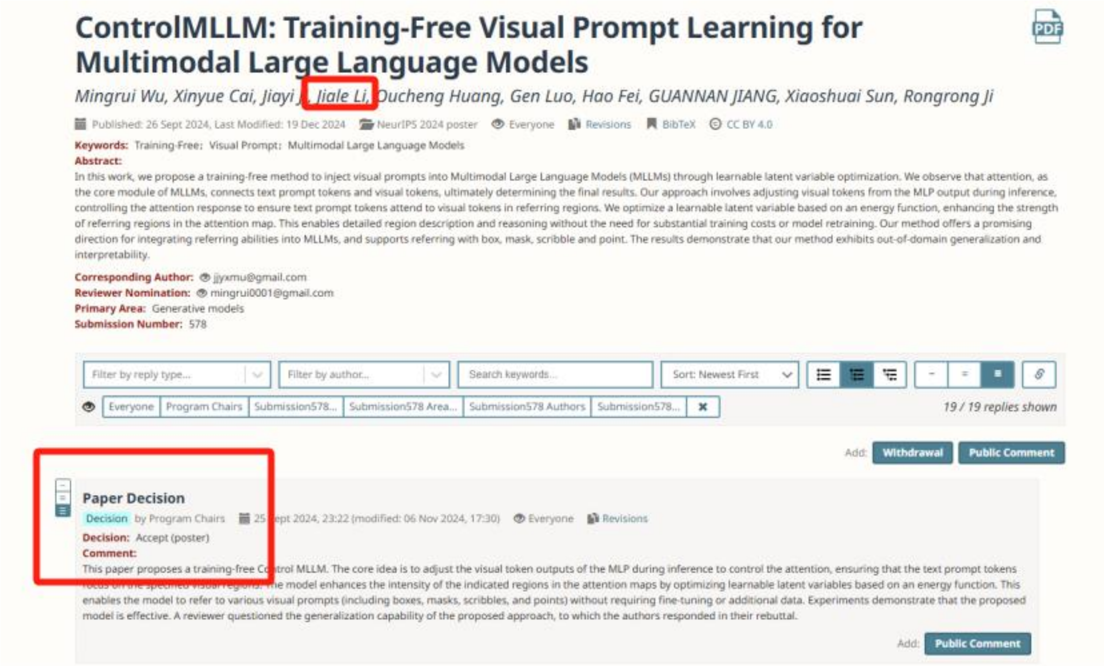

ControlMLLM: Training-Free Visual Prompt Learning for Multimodal Large Language Models

Mingrui Wu, Xinyue Cai, Jiayi Ji, Jiale Li, Oucheng Huang, Gen Luo, Hao Fei, GUANNAN JIANG, Xiaoshuai Sun, Rongrong Ji

Published: 26 Sept 2024, Last Modified: 19 Dec 2024 NeurIPS 2024 poster Everyone Revisions BibTeX CC BY 4.0

Keywords: Training-Free; Visual Prompt; Multimodal Large Language Models

Abstract:
In this work, we propose a training-free method to inject visual prompts into Multimodal Large Language Models (MLLMs) through learnable latent variable optimization. We observe that attention, as the core module of MLLMs, connects text prompt tokens and visual tokens, ultimately determining the final results. Our approach involves adjusting visual tokens from the MLP output during inference, controlling the attention response to ensure text prompt tokens attend to visual tokens in referring regions. We optimize a learnable latent variable based on an energy function, enhancing the strength of referring regions in the attention map. This enables detailed region description and reasoning without the need for substantial training costs or model retraining. Our method offers a promising direction for integrating referring abilities into MLLMs, and supports referring with box, mask, scribble and point. The results demonstrate that our method exhibits out-of-domain generalization and interpretability.

Corresponding Author: jjyxmu@gmail.com
Reviewer Nomination: mingrui0001@gmail.com
Primary Area: Generative models
Submission Number: 578

Filter by reply type... Filter by author... Search keywords... Sort: Newest First

Everyone Program Chairs Submission578... Submission578 Area... Submission578 Authors Submission578...

19 / 19 replies shown

Add: Withdrawal Public Comment

Paper Decision

Decision by Program Chairs 25 Sept 2024, 23:22 (modified: 06 Nov 2024, 17:30) Everyone Revisions

Decision: Accept (poster)

Comment:
This paper proposes a training-free Control MLLM. The core idea is to adjust the visual token outputs of the MLP during inference to control the attention, ensuring that the text prompt tokens focus on the specified visual regions. The model enhances the intensity of the indicated regions in the attention maps by optimizing learnable latent variables based on an energy function. This enables the model to refer to various visual prompts (including boxes, masks, scribbles, and points) without requiring fine-tuning or additional data. Experiments demonstrate that the proposed model is effective. A reviewer questioned the generalization capability of the proposed approach, to which the authors responded in their rebuttal.

Add: Public Comment

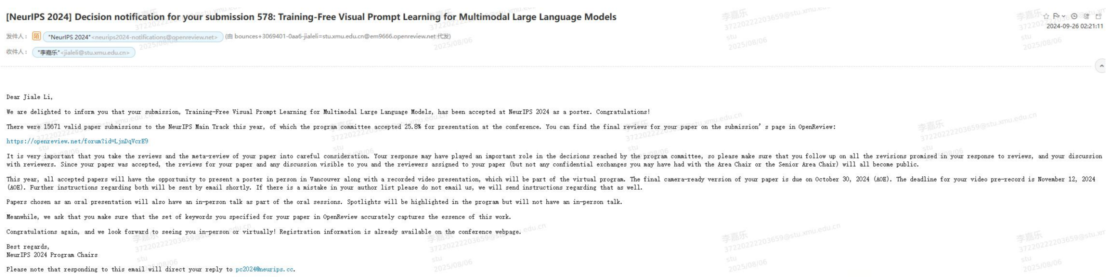

[NeurIPS 2024] Decision notification for your submission 578: Training-Free Visual Prompt Learning for Multimodal Large Language Models

2024-09-26 02:21:11

发件人： "NeurIPS 2024"<neurips2024-notifications@openreview.net> (由 bounces+3069401-0aa6-jialeli=stu.xmu.edu.cn@em9666.openreview.net 代发)

收件人： "李嘉乐"<jialeli@stu.xmu.edu.cn>

Dear Jiale Li,

We are delighted to inform you that your submission, Training-Free Visual Prompt Learning for Multimodal Large Language Models, has been accepted at NeurIPS 2024 as a poster. Congratulations!

There were 15671 valid paper submissions to the NeurIPS Main Track this year, of which the program committee accepted 25.8% for presentation at the conference. You can find the final reviews for your paper on the submission's page in OpenReview:

https://openreview.net/forum?id=LjnDqVcrE9

It is very important that you take the reviews and the meta-review of your paper into careful consideration. Your response may have played an important role in the decisions reached by the program committee, so please make sure that you follow up on all the revisions promised in your response to reviews, and your discussion with reviewers. Since your paper was accepted, the reviews for your paper and any discussion visible to you and the reviewers assigned to your paper (but not any confidential exchanges you may have had with the Area Chair or the Senior Area Chair) will all become public.

This year, all accepted papers will have the opportunity to present a poster in person in Vancouver along with a recorded video presentation, which will be part of the virtual program. The final camera-ready version of your paper is due on October 30, 2024 (AOE). The deadline for your video pre-record is November 12, 2024 (AOE). Further instructions regarding both will be sent by email shortly. If there is a mistake in your author list please do not email us, we will send instructions regarding that as well.

Papers chosen as an oral presentation will also have an in-person talk as part of the oral sessions. Spotlights will be highlighted in the program but will not have an in-person talk.

Meanwhile, we ask that you make sure that the set of keywords you specified for your paper in OpenReview accurately captures the essence of this work.

Congratulations again, and we look forward to seeing you in-person or virtually! Registration information is already available on the conference webpage.

Best regards,
NeurIPS 2024 Program Chairs

Please note that responding to this email will direct your reply to pc2024@neurips.cc.

ControlMLLM: Training-Free Visual Prompt Learning for Multimodal Large Language Models

Mingrui Wu[1], **Xinyue Cai**[1], **Jiayi Ji**[1]*, **Jiale Li**[1], **Oucheng Huang**[1],
Gen Luo[1], **Hao Fei**[2], **Guannan Jiang**[3], **Xiaoshuai Sun**[1], **Rongrong Ji**[1]
1 Key Laboratory of Multimedia Trusted Perception and Efficient Computing,
Ministry of Education of China, Xiamen University, 361005, P.R. China
2 National University of Singapore 3 CATL
mingrui0001@gmail.com

Abstract

In this work, we propose a training-free method to inject visual prompts into Multimodal Large Language Models (MLLMs) through learnable latent variable optimization. We observe that attention, as the core module of MLLMs, connects text prompt tokens and visual tokens, ultimately determining the final results. Our approach involves adjusting visual tokens from the MLP output during inference, controlling the attention response to ensure text prompt tokens attend to visual tokens in referring regions. We optimize a learnable latent variable based on an energy function, enhancing the strength of referring regions in the attention map. This enables detailed region description and reasoning without the need for substantial training costs or model retraining. Our method offers a promising direction for integrating referring abilities into MLLMs, and supports referring with box, mask, scribble and point. The results demonstrate that our method exhibits **out-of-domain generalization** and **interpretability**. Code: https://github.com/mrwu-mac/ControlMLLM.

1 Introduction

In recent times, there has been a surge in the development and adoption of large language models (LLMs), such as GPT-4 [1] and Llama [53], showcasing remarkable capabilities in addressing a wide range of human-generated questions. The success of these models has sparked interest among researchers in exploring the integration of LLMs with visual inputs. Consequently, a new class of models known as Multimodal Large Language Models (MLLMs) has emerged [36, 33, 17, 67, 81, 38]. However, despite their widespread adoption, traditional MLLMs often face limitations due to their reliance on coarse image-level alignments. This restricts users to guiding MLLMs solely through text prompts for detailed region description and reasoning. However, text often fails to capture the intricate visual nuances present in an image.

Addressing this challenge, recent efforts [68, 5, 75, 12, 37] have pioneered the integration of referring abilities within MLLMs, which enables users to provide input by pointing to specific coordinates of the objects or regions, as shown in Figure 1 (left). However, these endeavors typically entail substantial training costs to equip MLLMs with referring capabilities. Additionally, the model must undergo retraining to adapt to new data domains or new base MLLMs.

In this work, we propose a training-free method to inject the visual prompts into the Multimodal Large Language Models via learnable latent variable optimization. The method originates from our observation of the attention maps from the MLLM decoder, which model the relationship between

*Corresponding Author

38th Conference on Neural Information Processing Systems (NeurIPS 2024).

Evaluating and Analyzing Relationship Hallucinations in Large Vision-Language Models

Evaluating and Analyzing Relationship Hallucinations in Large Vision-Language Models

Mingrui Wu, Jiayi Ji, Oucheng Huang, Jiale Li, Yuhang Wu, Xiaoshuai Sun, Rongrong Ji

Published: 02 May 2024, Last Modified: 25 Jun 2024 ICML 2024 Poster Everyone Revisions BibTeX CC BY 4.0

Verify Author List: I have double-checked the author list and understand that additions and removals will not be allowed after the submission deadline.

Keywords: Relationship Hallucinations; Large Vision-Language Models

Abstract:
The issue of hallucinations is a prevalent concern in existing Large Vision-Language Models (LVLMs). Previous efforts have primarily focused on investigating object hallucinations, which can be easily alleviated by introducing object detectors. However, these efforts neglect hallucinations in inter-object relationships, which is essential for visual comprehension. In this work, we introduce R-Bench, a novel benchmark for evaluating Vision Relationship Hallucination. R-Bench features image-level questions that focus on the existence of relationships and instance-level questions that assess local visual comprehension. We identify three types of relationship co-occurrences that lead to hallucinations: relationship-relationship, subject-relationship, and relationship-object. The visual instruction tuning dataset's long-tail distribution significantly impacts LVLMs' understanding of visual relationships. Additionally, our analysis reveals that current LVLMs tend to overlook visual content, overly rely on the common sense knowledge of Large Language Models (LLMs), and struggle with spatial relationship reasoning based on contextual information.

Primary Area: Trustworthy Machine Learning (accountability, causality, fairness, privacy, robustness, etc.)

Position Paper Track: No

Paper Checklist Guidelines: I certify that all co-authors of this work have read and commit to adhering to the Paper Checklist Guidelines, Call for Papers and Publication Ethics.

Verify Author Names: My co-authors have confirmed that their names are spelled correctly both on OpenReview and in the camera-ready PDF. (If needed, please update 'Preferred Name' in OpenReview to match the PDF.)

No Additional Revisions: I understand that this submission is the final camera ready paper and that there will not be another opportunity to revise it. I have verified with all authors that they approve of this version.

Pdf Appendices: My camera-ready PDF file contains both the main text (not exceeding the page limits) and all appendices that I wish to include. I understand that any other supplementary material (e.g., separate files uploaded to OpenReview) will not be visible in the PMLR proceedings.

Latest Style File: I have compiled the camera ready paper with the latest ICML2024 style files (https://media.icml.cc/Conferences/ICML2024/Styles/icml2024.zip) and the compiled PDF includes an unnumbered Impact Statement section.

Paper Verification Code: YWQxM

Permissions Form: pdf

Supplementary Material: zip

Submission Number: 1112

Filter by reply type... Filter by author... Search keywords... Sort: Newest First

Everyone | Program Chairs | Submission1112 Authors | Submission1112... | Submission1112 Area... | Submission1112... | Submission1112... | Submission1112... | Submission1112...

21 / 21 replies shown

Add: **Withdrawal**

Paper Decision

Decision by Program Chairs 02 May 2024, 07:38 (modified: 06 Jun 2024, 23:47) Program Chairs, Authors Revisions

Decision: Accept (Poster)

Comment:
The paper explores relationship hallucinations in Large Vision-Language Models (LVLMs), introducing R-Bench, a novel benchmark specifically designed for evaluating such hallucinations. It focuses on image-level and instance-level questions to dissect LVLMs' understanding of inter-object relationships, identifying key issues like over-reliance on common sense and struggles with

[ICML 2024] Decision notification for your submission 1112: Evaluating and Analyzing Relationship Hallucinations in Large Vision-Language Models

2024-05-02 08:43:52

发件人： "ICML 2024"<icml2024-notifications@openreview.net> (由 bounces+3069401-0aa6-jialeli=stu.xmu.edu.cn@em9666.openreview.net 代发)

收件人： "李嘉乐"<jialeli@stu.xmu.edu.cn>

Dear Jiale Li,

We are happy to notify you that your submission Evaluating and Analyzing Relationship Hallucinations in Large Vision-Language Models (https://openreview.net/forum?id=xpSlt67vxQ), is accepted at ICML 2024. To access your reviews, please log in to your author console at https://openreview.net/group?id=ICML.cc/2024/Conference/Authors.

This year, ICML received 9,473 submissions (not including desk rejected papers), an increase of 44% from last year. Among these, we have accepted 2,609 submissions for presentation at the conference, an acceptance rate of 27.5%. These numbers include 286 position paper submissions, out of which 75 were accepted. Submissions were reviewed by at least three reviewers, an area chair, and a senior area chair, in order to ensure each submission was assessed properly; in the rare cases where a paper ultimately received only two reviews, ACs provided a detailed reading of the paper as well.

Note that while some of the meta-reviews mention recommending the paper for a talk versus a poster, **no final decisions have yet been made on the assignment of papers to presentations as a poster alone or as an oral presentation.** These will be made in the coming weeks, subject to availability of space and timing at the venue. All accepted papers will receive an email in the next few weeks designating the paper as a poster or additional oral presentation.

ICML will take place in Vienna, Austria on July 21-27 and will be an **in-person conference.** To accommodate this format, here is some important information:

1. Every paper will be presented as a poster at one of the poster sessions during the main conference.
2. Every paper will be given an opportunity to record and make available a short video presentation of the paper. This is encouraged but not required.
3. Papers eventually designated for oral presentation will receive further instructions on the presentation format.
4. Please update your paper as appropriate based on the review process. The camera ready deadline is May 29 (11:59pm AoE). Additional instructions and camera-ready requirements will be provided soon.

Congratulations, and thank you for submitting your work to ICML. We are looking forward to seeing you at the conference!

Sincerely,
ICML 2024 Program Chairs

Please note that responding to this email will direct your reply to program-chairs@icml.cc.

Evaluating and Analyzing Relationship Hallucinations in Large Vision-Language Models

Mingrui Wu [1] **Jiayi Ji** [1] **Oucheng Huang** **Jiale Li** [1] **Yuhang Wu** [1] **Xiaoshuai Sun** [1] **Rongrong Ji** [1]

Abstract

The issue of hallucinations is a prevalent concern in existing Large Vision-Language Models (LVLMs). Previous efforts have primarily focused on investigating object hallucinations, which can be easily alleviated by introducing object detectors. However, these efforts neglect hallucinations in inter-object relationships, which is essential for visual comprehension. In this work, we introduce R-Bench, a novel benchmark for evaluating Vision Relationship Hallucination. R-Bench features image-level questions that focus on the existence of relationships and instance-level questions that assess local visual comprehension. We identify three types of relationship co-occurrences that lead to hallucinations: relationship-relationship, subject-relationship, and relationship-object. The visual instruction tuning dataset's long-tail distribution significantly impacts LVLMs' understanding of visual relationships. Furthermore, our analysis reveals that current LVLMs tend to disregard visual content and overly rely on the common sense knowledge of Large Language Models. They also struggle with reasoning about spatial relationships based on contextual information.

1. Introduction

Recently, large language models (LLMs) such as GPT-4 (OpenAI, 2023) and Llama (Touvron et al., 2023) have demonstrated significant capabilities in addressing a broad spectrum of human-generated questions. The success of these models has spurred researchers to explore the use of LLMs in conjunction with visual inputs, leading to the development of various large vision-language models (LVLMs) (Li et al., 2023c; Liu et al., 2023c; Dai et al., 2023; Ye et al., 2023a). These endeavors typically involve methods like visual language pretraining (Li et al., 2023c) or visual instruction tuning (Liu et al., 2023c), aimed at integrating pre-trained visual encoders with LLMs to enhance their understanding of visual contexts.

However, despite their impressive performance, a significant challenge for these models is the unavoidable issue of hallucinations. Existing LVLMs often tend to generate responses that are inconsistent with the content of the images. This issue is particularly critical for LVLMs, which are expected to accurately comprehend images and produce answers consistent with the content of the visual input. While prior research has delved into evaluating object hallucinations (Li et al., 2023e), offering mitigation strategies through the object detection (Yin et al., 2023) or segmentation models (Wu et al., 2022; Chen et al., 2023c), there exists a notable gap in addressing hallucinations related to inter-object relationships, as shown in Figure 1. The latter, compared to object hallucinations, better reflects the LVLM's capacity to comprehend the intricacies of visual scenes. Currently, there is a shortage of comprehensive and rigorous benchmarks, to address these relationship hallucinations.

In this study, we introduce a novel Relationship Hallucination Benchmark (R-Bench) designed specifically for assessing relationship hallucinations in LVLMs. This benchmark comprises image-level and instance-level questions, labeled as 'Yes' or 'No', similar to the POPE evaluation (Li et al., 2023e). Image-level questions assess the existence of relationships in the image, while instance-level questions focus on specific object relationships, indicated by color bounding boxes or masks. The instance-level questions showcase the local visual understanding ability of LVLMs, adaptable to existing models without requiring retraining.

For the benchmark, we employ a combination of automatic generation by the Large Language Model (LLM) and manual curation. To ensure the benchmark's integrity, it is based on the nocaps validation set (Agrawal et al., 2019), preventing overlap with the pre-trained data of LVLMs. The construction process involves parsing all COCO captions to create a comprehensive relationship set. For each image in the nocaps dataset, we parse the provided captions

[1]Key Laboratory of Multimedia Trusted Perception and Efficient Computing, Ministry of Education of China, Xiamen University, 361005, P.R. China. Correspondence to: Jiayi Ji <jjyxmu@gmail.com>.

Proceedings of the 41st International Conference on Machine Learning, Vienna, Austria. PMLR 235, 2024. Copyright 2024 by the author(s).

基于图像识别的视频中服装自动检索与电商平台链接生成系统研究（已结项）

厦门大学“大学生创新训练计划”
创新训练项目申报书

学院：	信息学院
项目名称：	基于图像识别的视频中服装自动检索与电商平台链接生成系统研究
负责人：	李嘉乐
联系方式：	18563286057
指导教师：	晁飞
职称：	副教授
学历：	博士研究生毕业
E-mail：	1659589845@qq.com
项目申请日期：	2023年12月5日
项目起止年月：	2023年12月 至 2025年4月

项目名称	项目类型	年度	负责人	导师	负责学院	项目状态	刷新
基于图像识别的视频中服装自动检索与电商平台链接生成系统研究	创新训练	2024(1)	李嘉乐	晁飞	信息学院	结题已完成	[查看] [打印]

共 1 条 < 1 >

志愿工时

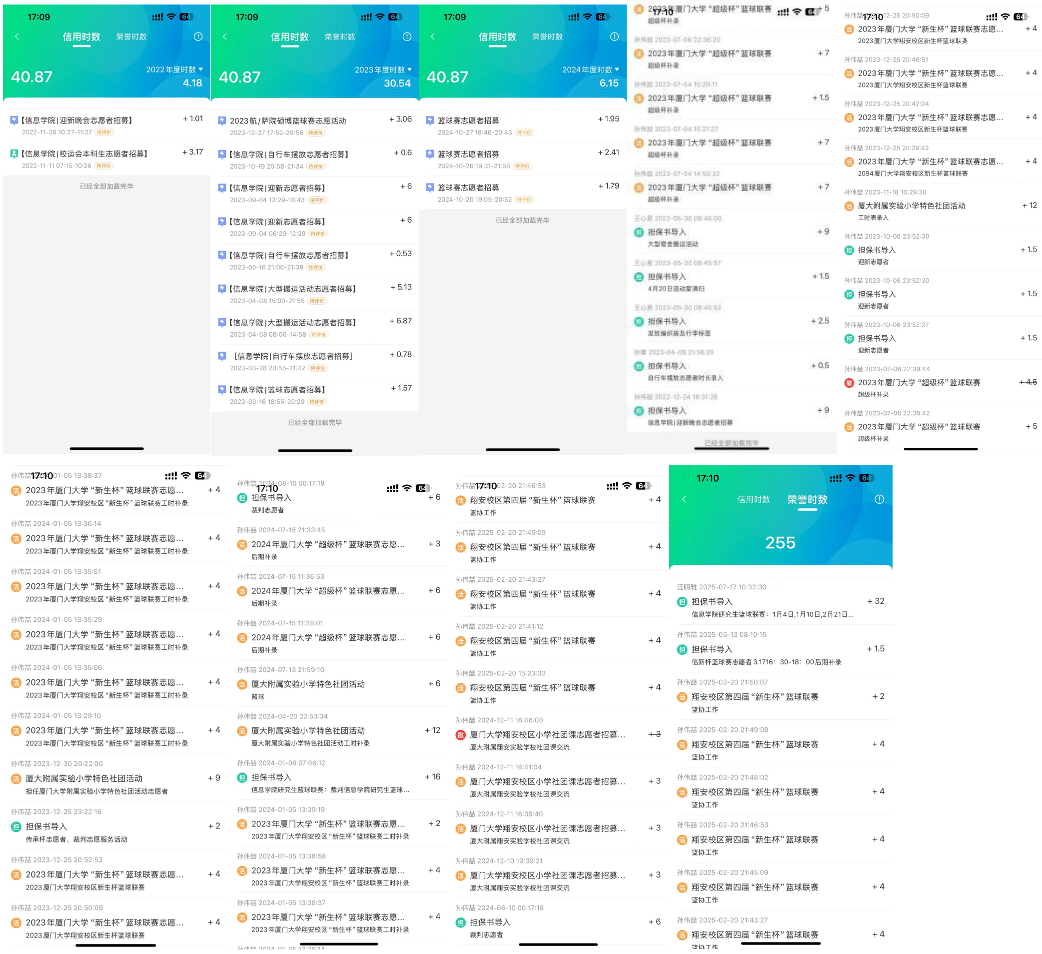

信用时数
荣誉时数
40.87
2022年度时数
4.18
2023年度时数
30.54
2024年度时数
6.15
255

两年班长

李嘉乐
2022级 信息学院
扫码查看“第二课堂成绩单”

社会工作
智能1班 (班级) | 班长
任职评分：99/100
任职时间：2024-10-01至2025-06-23
智能1班 (班级) | 班长
任职评分：97/100
任职时间：2023-10-01至2024-06-23

四六级

全国大学英语四级考试
成绩报告单

CET®

姓　　名：李嘉乐
学　　校：厦门大学
院　　系：信息学院
身份证号：370405200405142212

笔　试

准考证号：352012231105104
考试时间：2023年6月

总分	听力 (35%)	阅读 (35%)	写作和翻译 (30%)
524	150	202	172

口　试

准考证号：--
考试时间：--

成绩	--

成绩报告单编号：231135201004193

教育部教育考试院 证书专用章 11010210047228

校验码：FO2Y OYMZ T4X2 GDPY

说　明

1. 全国大学英语四、六级考试（CET）是由教育部主办的全国统一考试，考试对象为在校大学生。考试内容包括听、说、读、写、译等语言技能。
2. CET笔试考试时间为每年6月和12月；CET口试考试时间为每年5月和11月。
3. 考生可登录中国教育考试网（www.neea.edu.cn）查询、下载电子成绩报告单或自行办理纸质成绩证明。电子成绩报告单和纸质成绩证明与纸质成绩报告单具有同等效力。

大学英语四级口语考试能力描述

优秀	能用英语就熟悉的话题进行有效的交流；能清晰地叙述或描述一般性事件和现象。语言表达清楚连贯。
良好	能用英语就熟悉的话题进行交流；能叙述或描述一般性事件和现象。语言表达基本准确。
合格	能用英语就熟悉的话题进行简单交流；能简单叙述或描述一般性事件和现象。

全国大学英语六级考试
成绩报告单

CET®

姓　　名：李嘉乐
学　　校：厦门大学
院　　系：信息学院
身份证号：370405200405142212

笔　试

准考证号：352012242208308
考试时间：2024年12月

总分	听力 (35%)	阅读 (35%)	写作和翻译 (30%)
452	103	200	149

口　试

准考证号：--
考试时间：--

成绩	--

成绩报告单编号：242235201006606

教育部教育考试院 证书专用章 11010210047228

校验码：MEMX 2WCH DU2V J67W

说　明

1. 全国大学英语四、六级考试（CET）是由教育部主办的全国统一考试，考试对象为在校大学生。考试内容包括听、说、读、写、译等语言技能。
2. CET笔试考试时间为每年6月和12月；CET口试考试时间为每年5月和11月。
3. 考生可登录中国教育考试网（www.neea.edu.cn）查询、下载电子成绩报告单或自行办理纸质成绩证明。电子成绩报告单和纸质成绩证明与纸质成绩报告单具有同等效力。

大学英语六级口语考试能力描述

优秀	能用英语就一般性话题清晰地阐述自己的观点，明确地表达自己的态度；能开展深入的讨论，发表具有一定深度的见解。语言表达适切，自然流畅。
良好	能用英语就一般性话题阐述自己的观点，表明自己的态度；能开展较深入的讨论。语言表达准确连贯。
合格	能用英语就一般性话题进行交流；能参与讨论。语言表达基本准确。